

# Temporal-Difference Error Backpropagation Learning in Bidirectionally-Connected Neocortical Spiking Networks

Randall C. O'Reilly  
 Astera Institute Obelisk Lab  
 Departments of Psychology and Computer Science  
 Center for Neuroscience  
 University of California, Davis  
 1544 Newton Ct  
 Davis, CA 95616  
 oreilly@ucdavis.edu

December 9, 2022

Abstract:

Now that's abstract..

How does the human brain learn from experience? This basic question remains unanswered, specifically with respect to the most important brain system, the neocortex, which is responsible for all of our higher cognitive functions that are acquired over a lifetime of learning. The most powerful and well-established computational-level model is error backpropagation (Rumelhart, Hinton, & Williams, 1986), which drives the current generation of deep learning artificial intelligence (AI) systems (LeCun, Bengio, & Hinton, 2015). There are a number of existing ideas about how this form of learning could be supported by the neocortex (Lillicrap, Santoro, Marris, Akerman, & Hinton, 2020; Whittington & Bogacz, 2019). Here, we further develop one of these approaches, based on the *GeneRec* (generalized recirculation) approximation originally developed in O'Reilly (1996) (see Scellier & Bengio, 2017; Xie & Seung, 2003 for other closely related derivations), where the error signal is represented as a difference over time (i.e., a *temporal difference*) between an earlier *prediction* phase and a subsequent *outcome* phase. Each of these phases represents a coherent state of activity in the cortical network driven by the natural bidirectional propagation of signals through the ubiquitous bidirectional connectivity in the neocortex, without any need for specialized neuron types that explicitly compute and propagate an error signal. The detailed properties of the corticothalamic pathways naturally generate a temporal sequence of prediction-outcome states, supporting a form of predictive error-driven learning that provides an essentially unlimited supply of training signals to drive learning (?).

The existing biologically-based models of backpropagation-like learning use simplified continuous-valued rate-coded neural signals. In this paper, we show that the temporal-difference error-driven learning mechanism also works well using discrete spiking neurons, computed with standard conductance-based equations that accurately capture neocortical pyramidal neuron spiking dynamics (Brette & Gerstner, 2005; Gerstner & Naud, 2009). One critical addition to these equations was required to finally achieve some initial progress in learning, after many years of unsuccessful attempts: the addition of longer time-scale NMDA

---

Draft Manuscript: Do not cite or quote without permission.

R. C. O'Reilly is Chief Scientist at the Astera Institute Obelisk Lab, which may derive indirect benefit from the work presented here.

Supported by: ONR grants N00014-19-1-2684/ N00014-18-1-2116 (Deep learn), N00014-14-1-0670 / N00014-16-1-2128 (Bidir vis)

and GABA-B channels that previously have been used to simulate sustained working-memory firing in the prefrontal cortex (PFC) (Sanders, Berends, Major, Goldman, & Lisman, 2013). These channels have time constants of around 100 and 50 msec, respectively, and work synergistically together to stabilize a sparse distributed pattern of neural firing for a sufficient period of time to support the temporal-difference signals, which evolve over a period of around 200 msec in the model (?). These channels exist throughout the neocortex, but likely have a weaker influence in the posterior neocortex relative to the PFC, consistent with the more transient activity states in posterior vs. PFC areas (Miller, Erickson, & Desimone, 1996, e.g.). Various additional biologically-based mechanisms were also required to achieve fully robust large-scale learning, as detailed below.

There are a number of widely-discussed benefits of the discrete spiking (action potentials) in biological neurons, including efficiency from sparse communication, the ability to rapidly communicate an initial first-pass interpretation of a sensory input based on the timing of the first neurons to spike (Thorpe, Delorme, & Van Rullen, 2001), the ability to stochastically represent and sample from full probability distributions (as in variational neural networks) (McKee, Crandell, Chaudhuri, & O'Reilly, 2021; ?), and the enhanced ability to encode multiple signals over time relative to a single continuous rate-code signal. Thus, being able to drive powerful error-driven learning using discrete spiking neurons opens up many new possibilities relative to existing models. There are a number of ways that have been developed to somehow obtain the benefits of spiking within the context of error-backpropagation learning, such as converting rate-code models to spiking after training (cites), or directly using biologically-implausible gradient descent backpropagation with approximate derivative functions for the spiking activation dynamics (cites). However, none of these other approaches withstand scrutiny as a fully biologically plausible model of how the neocortex could actually perform backpropagation-like learning using spiking activation dynamics.

Many of the models that start with a more solid neuroscience-based foundation have focused on the *spike timing dependent plasticity (STDP)* paradigm, based on the striking finding of asymmetric patterns of synaptic increase (long term potentiation, LTP), vs decrease (long term depression, LTD) as a function of the temporal relationship between pre-synaptic vs. post-synaptic firing (Bi & Poo, 1998; ?). Despite the replicability of the very specific original STDP protocol (requiring 1 s pauses between pairs of otherwise isolated pre-post spikes), further variations of this protocol involving more than a pair of spikes and different patterns of overall timing have shown that STDP is better understood as a kind of artifact of this specific protocol, rather than a robust principle for understanding how neurons actually learn in the context of more naturalistic patterns of activity (Graupner & Brunel, 2012; Shouval, Wang, & Wittenberg, 2010). Furthermore, at a computational level, STDP, like many other similar forms of Hebbian learning, lacks a well-defined overall objective function that it is known to be minimizing through the local effects of learning — it is essentially just a heuristic idea that has no particular guarantee of producing useful overall learning outcomes. Although various modelers have managed to coax STDP-based models into doing some useful things (Kheradpisheh, Ganjtabesh, Thorpe, & Masquelier, 2018), there are no demonstrations that it reliably solves a wide variety of challenging computational problems by organizing representations across multiple layers.

By contrast, the computational power of error-backpropagation has been amply demonstrated in the rapid progress of deep learning models in recent years — it should be clear by now that it is the single most universally powerful such learning mechanism known, and thus it seems uncontroversial that it would be highly beneficial if the brain were somehow able to leverage the power of this form of learning. The framework presented here shows how a number of established neural mechanisms and channels, including the NMDA and GABA-B mentioned above, along with a number of others, can result in a powerful and scalable form of error backpropagation without committing *any* “errors of commission” — i.e., we are prepared to defend every single aspect of the model against the available neuroscience data, and welcome others to challenge anything present in the model.

Of particular importance is new, direct empirical evidence showing that pyramidal neuron synapses exhibit synaptic plasticity consistent with the temporal-difference form of error-driven learning (?), as predicted by the GeneRec and related models (O'Reilly, 1996). In this experiment, neurons were driven pre and postsynaptically with temporal patterns of activity organized into an initial *minus* phase prediction followed by a *plus* phase outcome (e.g., 25 Hz then 50 Hz) over a 200 msec window, while monitoring for changes in synaptic efficacy. The results showed all of the qualitative patterns for the temporal difference mechanism: LTP for increasing minus-plus, LTD for decreasing, and no change for consistent activity across a 200 msec theta cycle. These results hold despite differences in overall firing rates (e.g., stable 25 Hz and 50 Hz both showed no change on average).

The first section below provides a comprehensive overview of the model, followed by a detailed treatment of all the major equations used. Then, a sequence of increasingly complex previously-published models are tested using the new spiking-based error-driven learning, testing the relative learning properties in comparison to existing rate-code models, and other approaches to biologically-plausible backpropagation. Then, we explore models that specifically demonstrate the benefits of spiking-based activations for encoding probability distributions and more graded signals. Finally, we discuss this model in comparison to other approaches to training spiking models, and other relevant models and data, in the general discussion.

## Overview of the Axon Algorithm

### *Functional Advantages of Spikes*

Aside from biological fidelity, does discrete spiking actually afford any computational / functional advantages over the kind of rate code activation used in Leabra and most abstract neural networks? Perhaps surprisingly, there isn't a particularly strong answer to this question in the broader field, and more work needs to be done to specifically test and document the differences below (we plan to create an *Axoff* package which is as identical as possible, just with rate codes instead of spikes, to help address this question here).

Meanwhile, based on extensive experience with Axon and Leabra, here are some likely advantages of spikes:

1. **Graded behavior with fast initial responding:** Spiking networks can exhibit both very fast initial communication of information *and* finely graded, proportional responses. In Leabra, if you turn down the activation gain factor, you can get finely graded responses, but the network is slow to respond as activation takes a while to build and propagate, *or* you can get fast initial responses with high gain, but then it tends to exhibit strongly bimodal, non-graded behavior (especially in the context of bidirectional attractor dynamics). Simon Thorpe and colleagues have emphasized this point about the significant information value carried in the timing of the first wave of spikes references (Thorpe et al., 1996). After this first wave, reasonable rates of subsequent spiking (max around 100 Hz or every 10 msec) can send a graded rate-code-like signal over time.
2. **Graded, permeable attractors:** In rate-code networks, each neuron is continuously broadcasting its graded activation value on every time step, creating a "wall of activation" that is relatively impermeable to new signals. By contrast, in spiking networks, there are always gaps in the signal, which can allow new information to more easily penetrate and shape the ongoing "conversation" among neurons. One manifestation of this difference is that Leabra models typically require significant amounts of decay between trials, to allow new inputs to shape the network response, while Axon models do not. This difference is particularly notable in large, deep networks.

3. **Stochastic sampling:** The thresholded nature of spiking creates natural variability where small differences in timing can be magnified chaos-like into larger differences in which neurons end up spiking (the first to spike generally inhibits other neurons), contributing to an observed high level of variability in spiking network responding approximated by the Poisson distribution. This is a "natural" form of stochasticity without requiring actual added random noise, and it may be that it ends up retaining more of the relevant underlying signal as a result. Specifically, attempts to add noise in Leabra never ended up improving network performance, presumably because it directly diluted the signal with noise. By contrast, the natural stochasticity in Axon networks appears to do a good job of probabilistically sampling noise distributions references(McKee et al, 2021), and still allows high levels of performance (although consistently perfect responding is generally not observed; nor is it in people or other animals).
4. **Time is a real additional dimension:** Spiking networks exhibit a number of additional time-dependent effects that are absent in rate-code models with their continuous communication, including synchrony, temporal summation, bursting vs. pausing, etc. A number of theories have speculated about the additional signaling capabilities of synchrony or bursting, e.g., as an additional attentional or binding factor. In Axon, we don't specifically build in any such mechanisms, and the relevant learning mechanisms required to leverage this additional dimension are not necessarily obviously present in the biology (see the <https://github.com/ccnlab/kinase/tree/main/sims/kinasekinase> model and discussion). Nevertheless it is likely that there are important emergent temporal dynamics, and the version of the learning rule that *does* afford at least some sensitivity to coincident neural firing does actually work better in practice, so more work needs to be done to understand these issues in the context of the Axon model.

### *Activation: AdEx Conductance-based Spiking*

Axon uses the full conductance-based AdEx (adapting exponential) discrete spiking model of Gerstner and colleagues ([https://en.wikipedia.org/wiki/Exponential\\_integrate\\_and\\_fire](https://en.wikipedia.org/wiki/Exponential_integrate_and_fire) Wikipedia on AdEx), using normalized units as shown here: <https://docs.google.com/spreadsheets/d/1jn-NcXY4-y3pOw6inFOgPYlaQodrGIjcsAWkiD9f1FQ/edit#base=defaults>. Adaptation is implemented in a more realistic manner compared to standard AdEx, using the M-type medium time-scale M-AHP (after hyperpolarizing) channel, and two long time scales of sodium-gated potassium channels: KNa references (Kaczmarek, 2013). Leabra implemented a closer rate-code approximation to AdEx.

AdEx elides the very fast sodium-potassium channels that drive the action potential spiking, as captured in the seminal references Hodgkin & Huxley (1952) (HH) equations, which are 4th order polynomials, and thus require a very fine-grained time scale below 1 msec and / or more computationally expensive integration methods. The qualitative properties of these dynamics are instead captured using an exponential function in AdEx, which can be updated at the time scale of 1 msec.

Despite this simplification, AdEx supports neurophysiologically-based conductance equations so that any number of standard channel types can be added to the model, each with their own conductance function. See <https://github.com/emeryf/axon/tree/master/channels> for a description of the channel types supported, and the code implementing them, including NMDA and GABA-B as described next.

**Longer-acting, bistable NMDA and GABA-B currents:** An essential step for enabling spiking neurons to form suitably stable, selective representations for learning was the inclusion of both NMDA and GABA-B channels, which are voltage dependent in a complementary manner as captured in the References-Sanders et al, 2013 model (which provided the basis for the implementation here). These channels have long time constants and the voltage dependence causes them to promote a bistable activation state, with a

smaller subset of neurons that have extra excitatory drive from the NMDA and avoid extra inhibition from GABA-B, while a majority of neurons have the opposite profile: extra inhibition from GABA-B and no additional excitation from NMDA. With stronger conductance levels, these channels can produce robust active maintenance dynamics characteristic of layer 3 in the prefrontal cortex (PFC), but for posterior cortex, we use lower values that produce a weaker, but still essential, form of bistability. Without these channels, neurons all just "take turns" firing at different points in time, and there is no sense in which a small subset are engaged to represent a specific input pattern -- that had been a blocking failure in all prior attempts to use spiking in Leabra models.

**Auto-normalized, relatively scaled Excitatory Conductances:** As in Leabra, the excitatory synaptic input conductance ( $G_e$  in the code, known as *net input* in artificial neural networks) is computed as an average, not a sum, over connections, based on normalized weight values, which are subject to scaling on a projection level to alter relative contributions. Automatic scaling is performed to compensate for differences in expected activity level in the different projections. See section on projection-scaling for details. All of this makes it much easier to create models of different sizes and configurations with minimal (though still non-zero) need for additional parameter tweaking.

### *Temporal and Spatial Dynamics of Dendritic Integration*

A key dividing line in biological realism of neural models concerns the inclusion of separate dynamics for dendrites versus the soma, with a considerable literature arguing that significant computational functionality arises from nonlinear dynamics in the dendritic integration process. AdEx is a single compartment "point neuron" model (soma only), and obviously there is a major tradeoff in computational cost associated with modeling dendritic dynamics within individual neurons in any detail. In Axon, we have taken a middle ground (as usual), by including a separate dendritic membrane potential  $V_{mDend}$  that better reflects the dynamics of depolarization in the dendrites, relative to the standard AdEx  $V_m$  which reflects full integration in the soma. Voltage-gated channels localized in the dendrites, including NMDA and GABA-B, are driven by this  $V_{mDend}$ , and doing so results in significantly better performance vs. using the somatic  $V_m$ .

Furthermore, synaptic inputs are integrated first by separate projections, and then integrated into the full somatic conductances, and thus it is possible to implement nonlinear interactions among the different dendritic branches where these different projections may be organized. This is done specifically in the MSN (medium spiny neurons) of the basal ganglia in the `pcore` algorithm, and in the PT (pyramidal tract, layer 5IB intrinsic bursting) neurons also implemented in `pcore`. As noted above, each projection is also subject to different scaling factors, which while still linear, is critical for enabling models to function properly (e.g., top-down projections must in general be significantly weaker than bottom-up projections, to keep the models from hallucinating). These ways of capturing dendritic dynamics probably capture a reasonable proportion of the relevant functional properties in the biology, but more work with direct comparisons with fully detailed compartmental models is necessary to understand these issues better. See the appendix-dendritic-dynamics Appendix: Dendritic Dynamics for more discussion.

### *Inhibitory Competition Function Simulating Effects of Interneurons*

The pyramidal cells of the neocortex that are the main target of axon models only send excitatory glutamatergic signals via positive-only discrete spiking communication, and are bidirectionally connected. With all this excitation, it is essential to have pooled inhibition to balance things out and prevent runaway excitatory feedback loops. Inhibitory competition provides many computational benefits for reducing the dimensionality of the neural representations (i.e., *sparse* distributed representations) and restricting learning to only a small subset of neurons, as discussed extensively in the <https://CompCogNeuro.org> Comp Cog Neuro textbook. It is likely that the combination of positive-only weights and spiking activations, along with

inhibitory competition, is *essential* for enabling axon to learn in large, deep networks, where more abstract, unconstrained algorithms like the Boltzmann machine fail to scale (paper TBD).

Inhibition is provided in the neocortex primarily by the fast-spiking parvalbumin positive (PV+) and slower-acting somatostatin positive (SST+) inhibitory interneurons in the cortex referencesCardin, 2018. Instead of explicitly simulating these neurons, a key simplification in Leabra that eliminated many difficult-to-tune parameters and made the models much more robust overall was the use of a summary inhibitory function. This function directly computes a pooled inhibitory conductance  $G_i$  as a function of the feedforward (FF) excitation coming into a Pool of neurons, along with feedback (FB) from the activity level within the pool. Fortuitously, this same <https://github.com/emera/axon/tree/master/fffbFFFB> function works well with spiking as well as rate code activations, but it has some biologically implausible properties, and also at a computational level requires multiple layer-level `for` loops that interfere with full parallelization of the code.

Thus, we are now using the <https://github.com/emera/axon/tree/master/fsfffbFS-FFFB> *fast & slow* FFFB function that more explicitly captures the contributions of the PV+ and SST+ interneurons, and is based directly on FF and FB spikes, without requiring access to the internal  $G_e$  and  $Act$  rate-code variables in each neuron. See above link for more info. This function works even better overall than the original FFFB, in addition to providing a much more direct mapping onto the underlying biology.

See the `examples/inhib` model (from the CCN textbook originally) for an exploration of the basic excitatory and inhibitory dynamics in these models, comparing interneurons with FS-FFFB.

### *Kinase-based, Trace-enabled Error-backpropagation Learning*

A defining feature of Leabra, and Axon, is that learning is **error driven**, using a *temporal difference* to represent the error signal as a difference in two states of activity over time: *minus* (prediction) then *plus* (outcome). This form of error-driven learning is biologically plausible, by virtue of the use of bidirectional connectivity to convey error signals throughout the network. Any temporal difference arising anywhere in the network can propagate differences throughout the network -- mathematically approximating error backpropagation error gradients references(O'Reilly, 1996).

In the original Leabra (and early Axon) formulation, a version of the simple, elegant contrastive hebbian learning (CHL) rule was used:

$$dW = (x^+y^+) - (x^-y^-), \quad (1)$$

where the weight change  $dW$  is proportional to *difference* of two hebbian-like products of sending ( $x$ ) and receiving ( $y$ ) activity, between the plus and minus phase.

Axon now uses a different formulation of error-driven learning, that accomplishes three major objectives relative to CHL:

1. Enabling greater sensitivity to small error gradients that can accumulate over time, by computing the error in part based on *linear* net-input terms instead of the highly non-linear activation terms in CHL (representing something like a spiking rate).
2. Supporting a temporally-extended *eligibility trace* factor that provides a biologically-plausible way of approximating the computationally-powerful backprop-through-time (BPTT) algorithm references(Bellec et al, 2020).
3. More directly connecting to the underlying biochemical mechanisms that drive synaptic changes, in terms of calcium-activated *kinases*, as explored in the more biophysically detailed <https://github.com/ccnlab/kinase/tree/main/sims/kinasekinase> model.

The detailed derivation of this *kinase trace* learning mechanism is provided below: appendix-kinase-trace-learning-rule-derivationAppendix: Kinase-Trace Learning Rule Derivation, and summarized here. The algorithm is presented first at an abstract mathematical level, and then in terms of the underlying biological mechanisms that actually implement it.

Also, as originally developed in Leabra, the <https://github.com/emertax/axon/blob/master/deepdeep> package implements the extra anatomically-motivated mechanisms for *predictive* error-driven learning <https://ccnlab.org/papers/OReillyRussinZolfagharEtAl21.pdf>OReilly et al., 2021, where the minus phase represents a prediction and the plus phase represents the actual outcome. Biologically, we hypothesize that the two pathways of connectivity into the Pulvinar nucleus of the thalamus convey a top-down prediction and a bottom-up ground-truth outcome, respectively, providing an abundant source of error signals without requiring an explicit teacher.

### Error Gradient and Credit Assignment

Mathematically, error-driven learning has two components: the **error gradient**, which reflects the contribution of the receiving neuron to the overall network error, and the **credit assignment** factor that determines how much *credit / blame* for this error gradient to assign to each sending neuron:

$$dW = \text{Error} * \text{Credit}$$

In the simplest form of error-driven learning, the *delta rule*, these two terms are:

$$dW = (y^+ - y^-)x \quad (2)$$

where  $y^+$  is the *target* activity of the receiving neuron in the plus phase vs. its *actual* activity  $y^-$  in the minus phase (this difference representing the *Error* gradient), and  $x$  is the sending neuron activity, which serves as the credit assignment. Thus, more active senders get more of the credit / blame (and completely inactive neurons escape any).

In the mathematics of backpropagation, with a rearrangement of terms as in references(O'Reilly, 1996), the Error gradient factor in a bidirectionally-connected network can be computed as a function of the difference between *net input* like terms in the plus and minus phases, which are the dot product of sending activations times the weights:  $g = \sum_i w_i x_i$

as:

$$\text{Error} = (g^+ - g^-) \quad (3)$$

while the credit assignment factor is the sending activation times the derivative of the receiving activation:

$$\text{Credit} = x_i y' \quad (4)$$

### Error Gradient: Linear Net-Input and Nonlinear Activation

The actual form of error-gradient term used in Axon includes a contribution of the receiving unit activity in addition to the net-input term:

$$\text{Error} = (g^+ + \gamma y^+) - (g^- + \gamma y^-) \quad (5)$$

where  $\gamma$  (about .3 in effect by default) weights the contribution of receiving activity to the error term, relative to the net input.

The inclusion of the receiving activation, in deviation from the standard error backpropagation equations, is necessary for learning to work effectively in the context of inhibitory competition and sparse distributed

representations, where neurons have a highly nonlinear activation dynamic, and most are not active at all, despite having significant excitatory net input. In particular, it is often the case that a given neuron will be out-competed (inhibited) in the plus phase by other neurons, despite having relatively consistent levels of excitatory input across both phases. By including the activations in the error signal factor, this change in receiving activity level will show up in the error term, and cause the weight changes to reflect the fact that this neuron may not be as useful for the current trial. On the other hand, using purely activation-based terms in the error signal as in the CHL version above makes the learning *too* nonlinear and loses the beneficial graded nature of the linear net-input based gradient, where small error gradients can accumulate over time to continually drive learning. In models using CHL, learning would often just get "stuck" at a given intermediate level. Thus, as is often the case, a balanced approach combining both factors works best.

It is notable that the standard backpropagation equations do *not* include a contribution of the receiving activity in the error signal factor, and thus drive every unit to learn based on linearly re-weighted versions of the same overall error gradient signal. However, a critical difference between backprop and Axon is that backprop nets rely extensively on negative synaptic weight values, which can thus change the sign of the error-gradient factor. By contrast, the positive-only weights in Axon mean that the net-input factor is monotonically and positively related to the strength of the sending activations, resulting in much more homogenous error signals across the layer. Thus, the inclusion of the activation term in the error-signal computation can also be seen as a way of producing greater neural differentiation in the context of this positive-only weight constraint.

*Biological basis* From a biological perspective, the combination of net-input and receiving activity terms reflects the two main sources of Ca in the receiving neuron's dendritic spines: NMDA and VGCCs (voltage-gated calcium channels). NMDA channels are driven by sending-neuron glutamate release as in the net-input factor, and VGCCs are driven exclusively by receiving neuron spiking activity. Although the NMDA channel also reflects receiving neuron membrane depolarization due to the need for Mg (magnesium) ion unblocking, in practice there is often sufficient depolarization in the dendritic compartments, such that Ca influx is mostly a function of sending activity. Furthermore, the receiving depolarization factor in NMDA is also captured by the receiving activity term, which thus reflects both VGCC and the receiving activity dependent aspect of the NMDA channel.

In the Axon implementation, the NMDA and VGCC Ca influx is directly used, instead of the abstract equation above, providing an entirely biologically-grounded basis for the learning mechanism. This Ca influx is integrated over multiple cascaded variables that reflect the biochemical steps from Ca to calcium calmodulin (CaM) to CaMKII (CaM kinase II), which is a central kinase involved in controlling synaptic plasticity bidirectionally (LTP = long-term potentiation and LTD = long-term depression; referencesCoultrap & Bayer, 2012).

The critical subtraction necessary to extract the *Error* gradient from the temporal difference in this Ca signal over time is hypothesized to arise from the competitive binding dynamics between CaMKII and another kinase called DAPK1 (death-associated protein kinase 1), which has been shown to drive LTD when it out-competes CaMKII for binding to a particular location (N2B) on the NMDA receptor references(Goodell et al., 2017). Mathematically, if DAPK1 integrates at a slower rate than CaMKII, this competitive dynamic computes a temporal derivative exactly as required for the *Error* gradient equation. For example, when Ca increases in the later plus phase relative to the earlier minus phase, then the faster CaMKII responds more quickly, and thus out-competes DAPK1 and drives LTP (i.e., the temporal derivative is positive). Conversely, if Ca decreases in the plus vs. minus phase, DAPK1 remains more activated than CaMKII, and LTD ensues. When Ca is consistent over time, the two play to a draw, and the synaptic weight remains the same. For more details, see the biophysically detailed <https://github.com/ccnlab/kinase/tree/main/sims/kinasekinase> model. This model also shows how a fully continuous-time mechanism with no artificial knowledge of the theta cycle can drive appropriate error-driven learning.



Critically, there is now direct evidence supporting the idea that pyramidal neuron synapses actually follow this temporal difference dynamic references(Jang et al., 2023). In this experiment, neurons were driven pre and postsynaptically with patterns of minus -- plus phase activity (e.g., 25 Hz then 50 Hz) and changes in synaptic efficacy recorded. The results showed all of the qualitative patterns for the temporal difference mechanism: LTP for increasing minus-plus, LTD for decreasing, and no change for consistent activity across a 200 msec theta cycle. These results hold despite differences in overall firing rates.

#### *Credit Assignment: Temporal Eligibility Trace*

The extra mathematical steps taken in references(O'Reilly, 1996) to get from backpropagation to the CHL algorithm end up eliminating the factorization of the learning rule into clear Error vs. Credit terms. While this produces a nice simple equation, it makes it essentially impossible to apply the results of referencesBellec et al., (2020), who showed that backprop-through-time can be approximated through the use of a credit-assignment term that serves as a kind of temporal *eligibility trace*, integrating sender-times-receiver activity over a window of prior time steps.

By adopting the above form of error gradient term, we can now also adopt this trace-based credit assignment term as:

$$Credit = \langle xy \rangle_t y' \quad (6)$$

where the angle-bracket expression indicates an exponential running-average time-integration of the sender and receiver activation product over time.

The most computationally-effective form of learning goes one step further in computing this credit-assignment trace factor, by integrating spike-driven activity traces (representing calcium currents) within a given theta cycle of activity, in addition to integrating across multiple such theta-cycle "trials" of activity for the eligibility trace. This is significantly more computationally-expensive, as it requires synapse-level integration of the calcium currents on the millisecond-level timescale (a highly optimized form of computation is used so updates occur only at the time of spiking, resulting in a roughly 2x increase in computational cost overall).

The final term in the credit assignment factor is the derivative of the receiving activation function,  $y'$ , which would be rather difficult to compute exactly for the actual AdEx spiking dynamics used in Axon. Fortunately, using the derivative of a sigmoid-shaped logistic function works very well in practice, and captures the essential functional logic for this derivative: learning should be maximized when the receiving neuron is in its most sensitive part of its activation function (i.e., when the derivative is the highest), and minimized when it is basically "pinned" against either the upper or lower extremes. Specifically the derivative of the logistic is:

$$y' = y(1 - y) \quad (7)$$

which is maximal at  $y = .5$  and zero at either 0 or 1. In Axon, this is computed using a time-integrated spike-driven Ca-like term (CaSpkD), with the max value across the layer used instead of the fixed 1 constant. In addition, it is useful to use an additional factor that reflects the normalized difference in receiving spiking across the minus and plus phase, which can be thought of as an empirical measure of the sensitivity of the receiving neuron to changes over time:

$$y' = y(1 - y) \frac{y^+ - y^-}{MAX(y^+, y^-)} \quad (8)$$

#### *Stabilization and Rescaling Mechanisms*

A collection of biologically-motivated mechanisms are used to provide a stronger "backbone" or "spine" for the otherwise somewhat "squishy" learning that emerges from the above error-driven learning mecha-

nisms, serving to stabilize learning over longer time scales, and prevent parasitic positive feedback loops that otherwise plague these bidirectionally-connected networks. These positive feedback loops emerge because the networks tend to settle into stable attractor states due to the bidirectional, generally symmetric connectivity, and there is a tendency for a few such states to get broader and broader, capturing more and more of the "representational space". The credit assignment process, which is based on activation, contributes to this "rich get richer" dynamic where the most active neurons experience the greatest weight changes. We colloquially refer to this as the "hog unit" problem, where a small number of units start to hog the representational space, and it represents a major practical barrier to effective learning if not managed properly. Note that this problem does not arise in the vast majority of purely feedforward networks used in the broader neural network field, which do not exhibit attractor dynamics. However, this kind of phenomenon is problematic in other frameworks with the potential for such positive feedback loops, such as on-policy reinforcement learning or generative adversarial networks.

Metaphorically, various forms of equalizing taxation and wealth redistribution are required to level the playing field. The set of stabilizing, anti-hog mechanisms in Axon include:

1. **SWt:** structural, slowly-adapting weights. In addition to the usual learning weights driven by the above equations, we introduce a much more slowly-adapting, multiplicative  $SW_t$  that represents the biological properties of the dendritic *spine* -- these SWts "literally" give the model a spine! Spines are structural complexes where all the synaptic machinery is organized, and they slowly grow and shrink via genetically-controlled, activity-dependent protein remodeling processes, primarily involving the *actin* fibers also found in muscles. A significant amount of spine remodeling takes place during sleep -- so the SWt updating represents a simple model of sleep effects.

The SWt is multiplicative in the sense that larger vs. smaller spines provide more or less room for the AMPA receptors that constitute the adaptive weight value. The net effect is that the more rapid trial-by-trial weight changes are constrained by this more slowly-adapting multiplicative factor, preventing more extreme changes. Furthermore, the SWt values are constrained by a zero-sum dynamic relative to the set of receiving connections into a given neuron, preventing the neuron from increasing all of its weights higher and hogging the space. The SWt is also initialized with all of the randomness associated with the initial weights, and preserving this source of random variation, preventing weights from becoming too self-similar, is another important function.

2. **Target activity levels:** There is extensive evidence from Gina Turrigiano and collaborators, among others, that synapses are homeostatically rescaled to maintain target levels of overall activity, which vary across individual neurons (e.g., referencesTorrado Pacheco et al., 2021). Axon simulates this process, at the same slower timescale as updating the SWts (likewise associated with sleep), which are also involved in the rescaling process. The target activity levels can also slowly adapt over time, similar to an adaptive bias weight that absorbs the "DC" component of the learning signal referencesS-chraudolph (1998), but this adaptation is typically subject to a zero-sum constraint, so any increase in activity in one neuron must be compensated for by reductions elsewhere.

This is similar to a major function performed by the BCM learning algorithm in the Leabra framework -- by moving this mechanism into a longer time-scale outer-loop mechanism (consistent with Turrigiano's data), it worked much more effectively. By contrast, the BCM learning ended up interfering with the error-driven learning signal, and required relatively quick time-constants to adapt responsively as a neuron's activity started to change.

3. **Zero-sum weight changes:** In some cases it can also be useful to constrain the faster error-driven weight changes to be zero-sum, which is supported by an optional parameter. This zero-sum logic was

nicely articulated by referencesSchraudolph (1998), and is implemented in the widely-used ResNet models.

4. **Soft bounding and contrast enhancement:** To keep individual weight magnitudes bounded, we use a standard exponential-approach "soft bounding" dynamic (increases are multiplied by  $1-w$ ; decreases by  $w$ ). In addition, as developed in the Leabra model, it is useful to add a *contrast enhancement* mechanism to counteract the compressive effects of this soft bounding, so that effective weights span the full range of weight values.

## Axon Algorithm Equations

The pseudocode for Axon is given here, showing exactly how the pieces of the algorithm fit together, using the equations and variables from the actual code. Optimizations and special cases are omitted -- for that level of detail, the actual code is the definitive reference. For example, the time constants  $\tau$  are shown using division but in the code a pre-computed  $\Delta t = 1/\tau$  is used because multiplication is generally significantly faster than division.

timing-and-organization-of-computation

### Timing and Organization of Computation

Axon is organized around a 200 msec *theta* cycle (5 Hz), which is perhaps not coincidentally the modal peak for the duration of an eye fixation, and can be thought of as two 100 msec *alpha* cycles, which together comprise the minimal unit for predictive error driven learning according to the <https://github.com/emeric/axon/blob/master/deepdeep> predictive learning framework. Note that Leabra worked well with just a 100 msec alpha cycle, but it has not been possible to get the temporal difference error-driven learning mechanism to work at that scale with spiking in Axon, while it works very well at 200 msec. The time scales are:

- A **Cycle** represents 1 msec of processing, where each neuron is fully updated including all the conductances, integrated into the  $V_m$  membrane potential, which can then drive a *Spike* if it gets over threshold.
- A **Theta Cycle** or **Trial** or lasts 200 msec (200 Cycles; 5 Hz, theta frequency), and comprises one sequence of expectation -- outcome learning, with the final 50 msec comprising the *plus* phase when the outcome is active, while the preceding 150 msec is the *minus* phase when the network generates its own prediction or expectation. Ca influx from NMDA, VGCC are integrated continuously throughout this time, and drive learning at the end of the theta cycle. Ca at the synapse level is also integrated over time, reflecting intersection of pre and post activity, for the trace credit assignment factor in learning.

variables

### Variables

neuron

#### Neuron

The <https://github.com/emeric/axon/blob/master/axon/neuron.go> `axon.Neuron` struct contains all the neuron (unit) level variables, and the <https://github.com/emeric/axon/blob/master/axon/layer.go> `axon.Layer` contains a simple Go slice of these variables.

- *Spike* = whether neuron has spiked or not on this cycle (0 or 1)

- $Spiked = 1$  if neuron has spiked within the last 10 cycles (msecs), corresponding to a nominal max spiking rate of 100 Hz, 0 otherwise -- useful for visualization and computing activity levels in terms of average spiked levels
- $Act$  = rate-coded activation value reflecting instantaneous estimated rate of spiking, based on  $1 / ISI_{Avg}$ . This drives feedback inhibition in the FFFB function (todo: this will change when better inhibition is implemented), and is integrated over time for  $ActInt$  which is then used for performance statistics and layer average activations, etc. Should not be used for learning or other computations.
- $ActInt$  = integrated running-average activation value computed from  $Act$  to produce a longer-term integrated value reflecting the overall activation state across a reasonable time scale to reflect overall response of network to current input state -- this is copied to  $ActM$  and  $ActP$  at the ends of the minus and plus phases, respectively, and used in computing performance-level statistics (which are typically based on  $ActM$ ). Should not be used for learning or other computations.
- $ActM$  =  $ActInt$  activation state at end of third quarter, representing the posterior-cortical minus phase activation -- used for statistics and monitoring network performance. Should not be used for learning or other computations.
- $ActP$  =  $ActInt$  activation state at end of fourth quarter, representing the posterior-cortical plus phase activation -- used for statistics and monitoring network performance. Should not be used for learning or other computations.
- $Ext$  = external input: drives activation of unit from outside influences (e.g., sensory input).
- $Target$  = target value: drives learning to produce this activation value.
- $GeSyn$  = time-integrated total excitatory synaptic conductance, with an instantaneous rise time from each spike (in  $GeRaw$ ) and exponential decay with  $Dt.GeTau$ , aggregated over projections -- does *not* include  $Gbar.E$ .
- $Ge$  = total excitatory conductance, including all forms of excitation (e.g., NMDA) -- does *not* include  $Gbar.E$ .
- $GiSyn$  = time-integrated total inhibitory synaptic conductance, with an instantaneous rise time from each spike (in  $GiRaw$ ) and exponential decay with  $Dt.GiTau$ , aggregated over projections -- does *not* include  $Gbar.I$ . This is added with computed FFFB inhibition to get the full inhibition in  $Gi$ .
- $Gi$  = total inhibitory synaptic conductance -- the net inhibitory input to the neuron -- does *not* include  $Gbar.I$ .
- $Gk$  = total potassium conductance, typically reflecting sodium-gated potassium currents involved in adaptation effects -- does *not* include  $Gbar.K$ .
- $Inet$  = net current produced by all channels -- drives update of  $V_m$ .
- $V_m$  = membrane potential -- integrates  $Inet$  current over time.
- $V_mDend$  = dendritic membrane potential -- has a slower time constant, is not subject to the  $V_mR$  reset after spiking.

- $CaSyn$  = spike-driven calcium trace for synapse-level Ca-driven learning: exponential integration of  $SpikeG * Spike$  at  $SynTau$  time constant (typically 30). Synapses integrate  $send.CaSyn * recv.CaSyn$  across M, P, D time integrals for the synaptic trace driving credit assignment in learning. Time constant reflects binding time of Glu to NMDA and Ca buffering postsynaptically, and determines time window where pre \* post spiking must overlap to drive learning.
- $CaSpkM$  = spike-driven calcium trace used as a neuron-level proxy for synaptic credit assignment factor based on time-integrated spiking: exponential integration of  $SpikeG * Spike$  at  $MTau$  time constant (typically 5). Simulates a calmodulin (CaM) like signal at the most abstract level.
- $CaSpkP$  = cascaded integration of  $CaSpkM$  at  $PTau$  time constant (typically 40), representing neuron-level purely spiking version of plus, LTP direction of weight change and capturing the function of CaMKII in the Kinase learning rule. Used for specialized learning and computational functions, statistics, instead of Act.
- $CaSpkD$  = cascaded integration  $CaSpkP$  at  $DTau$  time constant (typically 40), representing neuron-level purely spiking version of minus, LTD direction of weight change and capturing the function of DAPK1 in the Kinase learning rule. Used for specialized learning and computational functions, statistics, instead of Act.
- $CaSpkPM$  = minus-phase snapshot of the  $CaSpkP$  value -- similar to  $ActM$  but using a more directly spike-integrated value.
- $CaLrn$  =  $recv$  neuron calcium signal used to drive temporal error difference component of standard learning rule, combining NMDA ( $NmdaCa$ ) and spiking-driven VGCC ( $VgccCaInt$ ) calcium sources (vs.  $CaSpk*$  which only reflects spiking component). This is integrated into  $CaM$ ,  $CaP$ ,  $CaD$ , and temporal derivative is  $CaP - CaD$  ( $CaMKII - DAPK1$ ). This approximates the backprop error derivative on net input, but VGCC component adds a proportion of  $recv$  activation delta as well -- a balance of both works best. The synaptic-level trace multiplier provides the credit assignment factor, reflecting coincident activity and potentially integrated over longer multi-trial timescales.
- $CaM$  = integrated  $CaLrn$  at  $MTau$  timescale (typically 5), simulating a calmodulin (CaM) like signal, which then drives  $CaP$ ,  $CaD$  for delta signal driving error-driven learning.
- $CaP$  = cascaded integration of  $CaM$  at  $PTau$  time constant (typically 40), representing the plus, LTP direction of weight change and capturing the function of CaMKII in the Kinase learning rule.
- $CaD$  = cascaded integratoin of  $CaP$  at  $DTau$  time constant (typically 40), representing the minus, LTD direction of weight change and capturing the function of DAPK1 in the Kinase learning rule.
- $CaDiff$  = difference between  $CaP - CaD$  -- this is the error signal that drives error-driven learning.
- $SpkMaxCa$  = Ca integrated like  $CaSpkP$  but only starting at  $MacCycStart$  cycle, to prevent inclusion of carryover spiking from prior theta cycle trial -- the  $PTau$  time constant otherwise results in significant carryover.
- $SpkMax$  = maximum  $CaSpkP$  across one theta cycle time window -- used for specialized algorithms that have more phasic behavior within a single trial, e.g., BG Matrix layer gating. Also useful for visualization of peak activity of neurons.
- $SpkPrv$  = final  $CaSpkD$  activation state at end of previous theta cycle. used for specialized learning mechanisms that operate on delayed sending activations.

- `SpkSt1` = the activation state at specific time point within current state processing window (e.g., 50 msec for beta cycle within standard theta cycle), as saved by `SpkSt1()` function. Used for example in hippocampus for CA3, CA1 learning.
- `SpkSt2` = the activation state at specific time point within current state processing window (e.g., 100 msec for beta cycle within standard theta cycle), as saved by `SpkSt2()` function. Used for example in hippocampus for CA3, CA1 learning.
- `RLRate` = recv-unit based learning rate multiplier, reflecting the sigmoid derivative computed from the `CaSpkD` of recv unit, and the normalized difference  $\text{CaSpkP} - \text{CaSpkD} / \text{MAX}(\text{CaSpkP} - \text{CaSpkD})$ .
- `ActAvg` = average activation (of minus phase activation state) over long time intervals (time constant = `Dt.LongAvgTau`) -- useful for finding hog units and seeing overall distribution of activation.
- `AvgPct` = `ActAvg` as a proportion of overall layer activation -- this is used for synaptic scaling to match `TrgAvg` activation -- updated at `SlowInterval` intervals
- `TrgAvg` = neuron's target average activation as a proportion of overall layer activation, assigned during weight initialization, driving synaptic scaling relative to `AvgPct`.
- `DTrgAvg` = change in neuron's target average activation as a result of unit-wise error gradient -- acts like a bias weight. MPI needs to share these across processors.
- `AvgDif` = `AvgPct` - `TrgAvg` -- i.e., the error in overall activity level relative to set point for this neuron, which drives synaptic scaling -- updated at `SlowInterval` intervals.
- `Attn` = Attentional modulation factor, which can be set by special layers such as the TRC -- multiplies `Ge`.
- `ISI` = current inter-spike-interval -- counts up since last spike. Starts at -1 when initialized.
- `ISIAvg` = average inter-spike-interval -- average time interval between spikes, integrated with `ISITau` rate constant (relatively fast) to capture something close to an instantaneous spiking rate. Starts at -1 when initialized, and goes to -2 after first spike, and is only valid after the second spike post-initialization.
- `GeNoiseP` = accumulating poisson probability factor for driving excitatory noise spiking -- multiply times uniform random deviate at each time step, until it gets below the target threshold based on `lambda`.
- `GeNoise` = integrated noise excitatory conductance, added into `Ge`.
- `GiNoiseP` = accumulating poisson probability factor for driving inhibitory noise spiking -- multiply times uniform random deviate at each time step, until it gets below the target threshold based on `lambda`.
- `GiNoise` = integrated noise inhibitory conductance, added into `Gi`.
- `GeM` = time-averaged `Ge` value over the minus phase -- useful for stats to set strength of connections etc to get neurons into right range of overall excitatory drive.
- `GiM` = time-averaged `GiSyn` value over the minus phase -- useful for stats to set strength of connections etc to get neurons into right range of overall excitatory drive.

- `MahpN` = accumulating voltage-gated gating value for the medium time scale AHP.
- `SahpCa` = slowly accumulating calcium value that drives the slow AHP.
- `SahpN` = sAHP gating value.
- `GknaMed` = conductance of sodium-gated potassium channel (KNa) medium dynamics (Slick) -- produces accommodation / adaptation of firing.
- `GknaSlow` = conductance of sodium-gated potassium channel (KNa) slow dynamics (Slack) -- produces accommodation / adaptation of firing.
- `GnmdaSyn` = integrated NMDA recv synaptic current -- adds `GeRaw` and decays with time constant.
- `Gnmda` = net postsynaptic (recv) NMDA conductance, after Mg V-gating and `Gbar` -- added directly to `Ge` as it has the same reversal potential.
- `GnmdaLrn` = learning version of integrated NMDA recv synaptic current -- adds `GeRaw` and decays with time constant -- drives `NmdaCa` that then drives `CaM` for learning.
- `NmdaCa` = NMDA calcium computed from `GnmdaLrn`, drives learning via `CaM`.
- `SnmdaO` = Sender-based number of open NMDA channels based on spiking activity and consequent glutamate release for all sending synapses -- this is the presynaptic component of NMDA activation that can be used for computing Ca levels for learning -- increases by  $(1 - \text{SnmdaI}) * (1 - \text{SnmdaO})$  with spiking and decays otherwise.
- `SnmdaI` = Sender-based inhibitory factor on NMDA as a function of sending (presynaptic) spiking history, capturing the allosteric dynamics from Urakubo et al (2008) model. Increases to 1 with every spike, and decays back to 0 with its own longer decay rate.
- `GgabaB` = net GABA-B conductance, after `Vm` gating and `Gbar` + `Gbase` -- applies to `Gk`, not `Gi`, for GIRK, with .1 reversal potential.
- `GABAB` = GABA-B / GIRK activation -- time-integrated value with rise and decay time constants.
- `GABABx` = GABA-B / GIRK internal drive variable -- gets the raw activation and decays.
- `Gvgcc` = conductance (via Ca) for VGCC voltage gated calcium channels.
- `VgccM` = activation gate of VGCC channels.
- `VgccH` = inactivation gate of VGCC channels.
- `VgccCa` = instantaneous VGCC calcium flux -- can be driven by spiking or directly from `Gvgcc`.
- `VgccCaInt` = time-integrated VGCC calcium flux -- this is actually what drives learning.
- `GeExt` = extra excitatory conductance added to `Ge` -- from `Ext` input, `deep.GeCtxt` etc.
- `GeRaw` = raw excitatory conductance (net input) received from senders = current raw spiking drive.
- `GeBase` = baseline level of `Ge`, added to `GeRaw`, for intrinsic excitability.
- `GiRaw` = raw inhibitory conductance (net input) received from senders = current raw spiking drive.

- $G_{iBase}$  = baseline level of  $G_i$ , added to  $G_{iRaw}$ , for intrinsic excitability.
- $SSG_i$  = SST+ somatostatin positive slow spiking inhibition.
- $SSG_{iDend}$  = amount of SST+ somatostatin positive slow spiking inhibition applied to dendritic Vm ( $V_{mDend}$ ).
- $G_{ak}$  = conductance of A-type K potassium channels.

inhib-pools

### *Inhib Pools*

There is always at least one <https://github.com/emerg/axon/blob/master/axon/pool.go> `axon.Pool` pool of neurons over which inhibition is computed, and there can also be pools for subsets of neurons that correspond to hypercolumns (or max pooling in a convolutional neural network), and support more local inhibitory dynamics.

- $FFsRaw$  = all feedforward incoming spikes into neurons in this pool -- raw aggregation.
- $FBsRaw$  = all feedback outgoing spikes generated from neurons in this pool -- raw aggregation.
- $FFs$  = all feedforward incoming spikes into neurons in this pool, normalized by pool size.
- $FBs$  = all feedback outgoing spikes generated from neurons in this pool, normalized by pool size.
- $FS_i$  = fast spiking PV+ fast integration of FFs feedforward spikes.
- $SS_i$  = slow spiking SST+ integration of FBs feedback spikes.
- $SSf$  = slow spiking facilitation factor, representing facilitating effects of recent activity.
- $FSG_i$  = overall fast-spiking inhibitory conductance.
- $SSG_i$  = overall slow-spiking inhibitory conductance.
- $G_i$  = overall inhibitory conductance =  $FSG_i + SSG_i$ .

synapse

### *Synapse*

Neurons are connected via synapses parameterized with the following variables, contained in the <https://github.com/emerg/axon/blob/master/axon/synapse.go> `axon.Synapse` struct. The <https://github.com/emerg/axon/blob/master/axon/prjn.go> `axon.Prjn` contains all of the synaptic connections for all the neurons across a given layer -- there are no Neuron-level data structures in the Go version.

- $W_t$  = effective synaptic weight value, determining how much conductance one spike drives on the receiving neuron, representing the actual number of effective AMPA receptors in the synapse.  $W_t = SW_t * W_tSig(LW_t)$ , where  $W_tSig$  produces values between 0-2 based on  $LW_t$ , centered on 1.
- $LW_t$  = rapidly learning, linear weight value -- learns according to the `lrate` specified in the connection spec. Biologically, this represents the internal biochemical processes that drive the trafficking of AMPA receptors in the synaptic density. Initially all  $LW_t$  are .5, which gives 1 from  $W_tSig$  function.



- $SW_t$  = slowly adapting structural weight value, which acts as a multiplicative scaling factor on synaptic efficacy: biologically represents the physical size and efficacy of the dendritic spine.  $SW_t$  values adapt in an outer loop along with synaptic scaling, with constraints to prevent runaway positive feedback loops and maintain variance and further capacity to learn. Initial variance is all in  $SW_t$ , with  $LW_t$  set to .5, and scaling absorbs some of  $LW_t$  into  $SW_t$ .
- $DW_t$  = change in synaptic weight, from learning -- updates  $LW_t$  which then updates  $W_t$ .
- $DSW_t$  = change in  $SW_t$  slow synaptic weight -- accumulates  $DW_t$
- $Ca$  = Raw calcium signal for Kinase learning:  $SpikeG * (send.CaSyn * recv.CaSyn)$
- $CaM$  = first stage running average (mean) Ca calcium level (like  $CaM$  = calmodulin), feeds into  $CaP$
- $CaP$  = shorter timescale integrated  $CaM$  value, representing the plus, LTP direction of weight change and capturing the function of CaMKII in the Kinase learning rule
- $CaD$  = longer timescale integrated  $CaP$  value, representing the minus, LTD direction of weight change and capturing the function of DAPK1 in the Kinase learning rule
- $Tr$  = trace of synaptic activity over time -- used for credit assignment in learning.

activation-update-cycle-every-1-msec-ge-gi-vm-spike

### *Activation Update Cycle (every 1 msec): Ge, Gi, Vm, Spike*

The `axon.NetworkCycleImpl` method in <https://github.com/emer/axon/blob/master/axon/network.go#L100> calls the following functions in order:

- `GFmSpikes` on all `Prjns`: integrates Raw and Syn conductances for each `Prjn` from spikes sent previously, into `GVals` organized by receiving neuron index, so they can then be integrated into the full somatic conductances in `CycleNeuron`.
- `GiFmSpikes` on all `Layers`: computes inhibitory conductances based on total incoming FF and FB spikes into the layer, using the <https://github.com/emer/axon/tree/master/fsfffbFS-FFFB> summary functions.
- `CycleNeuron` on all `Neurons`: integrates the `Ge` and `Gi` conductances from above, updates all the other channel conductances as described in <https://github.com/emer/axon/tree/master/chanschans>, and then computes `Inet` as the net current from all these conductances, which then drives updates to `Vm` and `VmDend`. If `Vm` exceeds threshold then `Spike = 1`. It also updates the neuron-level calcium variables that drive learning (`CaLrn`, `CaM`, `CaP`, `CaD` and `CaSpk` versions of these).
- `SendSpike` on all `Neurons`: for each neuron with `Spike = 1`, adds scaled synaptic weight value to `GBuf` ring buffer for efficiently delaying receipt of the spike per parametrized `Com.Delay` cycles. This is what the `GFmSpikes` then integrates. This is very expensive computationally because it goes through synapses on a msec Cycle scale.
- `CyclePost` on all `Layers`: a hook for specialized algorithms to do something special.
- `SendSynCa` and `RecvSynCa` on all `Prjns`: update synapse-level calcium (`Ca`) for any neurons that spiked (on either the send or recv side). This is expensive computationally because it goes through synapses on each msec Cycle (but only for neurons that actually spiked).

All of the relevant parameters and most of the equations are in the <https://github.com/emerg/axon/blob/master/axon/act.go>, <https://github.com/emerg/axon/blob/master/axon/inhib.go>, and <https://github.com/emerg/axon/blob/master/axon/learn.go> which correspond to the Act, Inhib and Learn fields in the Layer struct. Default values of parameters are shown in comments below.

gfmspikes-for-each-prjn

*GFmSpikes: for each Prjn*

Prjn.GVals integrates two synaptic conductance values G per receiving neuron index, using Ge time constants for excitatory synapses, and Gi for inhibitory. These values were sent before in SendSpike and stored in the Prjn.GBuf slice (see below). In principle synaptic conductance is computed using a standard *alpha function* double-exponential with one time constant for the rise and another for the decay. In practice, the rise time is < 1 msec and thus it is simpler to just add the new raw and decay (decay tau = 5 msec):

- GRaw = GBuf at the Com.Delay index (e.g., 2 msec)
- GSyn += GRaw - Act.Dt.GeDt \* GSyn // GeSynFmRaw or GiSynFmRaw

gifmspikes-for-each-layer

*GiFmSpikes: for each Layer*

Layer.Pool[\*].Inhib pools are updated based on FFsRaw and FBsRaw which are accumulated during SendSpike

Normalize raw values:

- FFs = FFsRaw / Npool // Npool = number of neurons in pool
- FBs = FBsRaw / Npool

Fast spiking (FS) PV from FFs and FBs, with decay:

- FSi = FFs + FB \* FBs - FSi / FSTau // FB = feedback weight, 1;  
FSTau = decay, 6 msec
- FSGi = Gi \* (|FSi - FS0|>0) // |x|>0 = only above 0; FS0 = 0.1

Slow spiking (SS) SST from FBs only, with facilitation factor SSf:

- SSi += (SSf \* FBs - SSi) / SSiTau // SSiTau = 50 msec
- SSf += FBs \* (1 - SSf) - SSf / SSfTau // SSfTau = 20 msec
- SSGi = Gi \* SS \* SSi // Gi = overall scaling, 1; SS = extra SS factor, 30

Overall inhibition:

- Gi = FSGi + SSGi

cycleneuron

*CycleNeuron*

There are three major steps here: `GInteg`, `SpikeFmG`, and `CaFmSpike`, the last of which updates `Ca` values used in learning.

`ginteg-integrate-graw-and-gsyn-from-recv-prjns-other-gs`

Iterates over `Recv Prjns`:

- `GeRaw += prjn.GRaw // if excitatory`
- `GeSyn += prjn.Gsyn`
- or `Gi*` if inhibitory

Then all the special conductances:

- NMDA, VGCC, GABAB, `Gk` -- see <https://github.com/emeric/axon/tree/master/chanschans> for equations, which operate on `VmDend` instead of `Vm`, as these channels are primarily located in the dendrites. These contribute to overall `Ge` excitatory conductance and `Gi` inhibition.

And add in the pool inhib `Gi` computed above.

`spikefmg-compute-vm-and-spikes-from-all-the-gs`

`Vm` is incremented by the net current `Inet` summing all the conductances and an exponential factor capturing the Hodgkin Huxley spiking dynamics:

- `Inet = Gbar.E * Ge * (Erev.E - Vm) + Gbar.I * Gi * (Erev.I - Vm) + Gbar.L * (Erev.L - Vm) + Gbar.K * Gk * (Erev.K - Vm)`

- `// Gbar.E = 1, I = 1, L = 0.2, K = 1; Erev.E = 1, L = 0.3, I = 0.1, K = 0.1`

- See [https://docs.google.com/spreadsheets/d/1jn-NcXY4-y3pOw6inFOgPYlaQodrGIjcsAWkiD9f1FQ/edit?usp=](https://docs.google.com/spreadsheets/d/1jn-NcXY4-y3pOw6inFOgPYlaQodrGIjcsAWkiD9f1FQ/edit?usp=sharing) sheet for conversions from biological values to normalized units used in model.

- `Vm += (Inet + Gbar.L * ExpSlope * Exp((Vm-Thr) / ExpSlope)) / VmTau`

- `// VmTau = 2.81 (capacitance); ExpSlope = 0.02 (2 mV biological)`

In the Axon implementation, 2 smaller steps are taken in integrating the `Vm` (i.e., .5 msec resolution), and a midpoint value is used in computing the exponential factor, to avoid numerical instability while maintaining a 1 msec overall update rate.

If the neuron has just spiked within the `Tr` refractory time window (3 msec default), then `Vm` just decays toward a refractory potential `VmR` (0.3 = rest potential) with a time constant of `RTau` (1.6667 msec), with the last step constrained to reach `VmR` exactly.

$V_{mDend}$  is updated in the same way as  $V_m$ , except that the exponential term is muted by an additional factor of  $G_{barExp}$  (0.2), and it is still updated by conductances during the refractory period, with an additional leak conductance of  $G_{barR}$  (3) added to drive the potential downward toward the resting potential. Thus, consistent with detailed compartmental models and electrophysiological data,  $V_{mDend}$  exhibits more sustained depolarization, which keeps the NMDA and GABA-B currents more stably activated, as shown in the appendix-dendritic-dynamicsAppendix: Dendritic Dynamics.

$V_{mDend}$  also has an additional contribution from the  $SSGi$  slow-spiking inhibition ( $2 * SSGi$  by default), reflecting the fact that the SST+ neurons target the distal dendrites. This is important functionally to counter a positive feedback loop from NMDA channels, as discussed here: <https://github.com/emery/axon/tree/master/fsfffbFS-FFFB>.

Spike is set to 1 when  $V_m > ExpThr$  (0.9, for default case where Exp function is being used), and the ISI (inter-spike-interval) counter, and time-averaged  $ISIAvg$  are updated:

- if  $ISIAvg \leq 0$  then  $ISIAvg = ISI$ ; else if  $ISI < 0.8 * ISIAvg$  then  $ISIAvg = ISI$ ; else  $ISIAvg += (ISI - ISIAvg) / ISITau$

If the neuron did not spike, then  $ISI++$  is incremented.

cafmspike-calrn-nmda-vgcc-and-simple-spike-driven-ca-signals

The core Ca calcium value that drives the *trace - kinase* learning rule is stored in the  $CaLrn$  neuron variable, as a sum of NMDA and VGCC calcium influx:

- $CaLrn = (NmdaCa + VgccCaInt) / Norm$

Where  $Norm$  (80) renormalizes the concentration-based factors to a range that works well for learning.

In larger networks, directly using the calcium flux from the VGCC channel ( $VgccCa$ ) works well, but in smaller networks it typically works better to use a simpler approximation to the VGCC Ca that purely reflects spiking rate and not the other voltage-related factors that affect Ca in the actual channels. In either case, the  $VgccCa$  is integrated over time with a decay constant reflecting buffering and diffusion of these highly transient Ca signals driven by spiking:

[] if SpkVGCC {  $VgccCa = SpkVgccCa * Spike // SpkVgccCa = 35$  }  $VgccCaInt += VgccCa - VgccCaInt / VgccTau // VgccTau = 10 \text{ msec}$

This immediate  $CaLrn$  value is then subject to multiple levels of additional integration processes, reflecting the CaM calmodulin  $\rightarrow$  CaMKII  $\rightarrow$  DAPK1 cascades, into the CaM, CaP and CaD variables. The same time constants are used for these processes across various different variables, and are defined in the <https://github.com/emery/axon/tree/master/kinasekinase> package, as follows:

- $MTau$  (2 or 5 msec) for CaM or  $CaSpkM$  = calmodulin time constant in cycles (msec) -- for synaptic-level integration this integrates on top of Ca signal from send- $\rightarrow$ CaSyn \* recv- $\rightarrow$ CaSyn, each of which are typically integrated with a 30 msec Tau.
- $PTau$  (40 msec) for CaP or  $CaSpkP$  = LTP spike-driven Ca factor time constant in cycles (msec), simulating CaMKII in the Kinase framework, with 40 on top of  $MTau$  roughly tracking the biophysical rise time. Computationally, CaP represents the plus phase learning signal that reflects the most recent past information.

- $DTau$  (40 msec) for  $CaD$  or  $CaSpkD$  = LTD spike-driven Ca factor time constant in cycles (msec), simulating DAPK1 in Kinase framework. Computationally,  $CaD$  represents the minus phase learning signal that reflects the expectation representation prior to experiencing the outcome (in addition to the outcome).

The cascading update looks like this:

```
[] CaM += (CaLrn - CaM) / MTau  CaP += (CaM - CaP) / PTau  CaD += (CaP - CaD) / DTau
```

As shown below for the  $DWt$  function, the *Error* gradient is:

- $Error = CaP - CaD$

In addition, the Ca trace used for synaptic-level integration for the trace-based *Credit* assignment factor (see `SendSynCa` and `RecvSynCa` below) is updated with a time constant of  $SynTau$  (30 msec) which defines the time window over which pre \* post synaptic activity interacts in updating the synaptic-level trace:

- $CaSyn += (SpikeG * Spike - CaSyn) / SynTau$

Finally, various peripheral aspects of learning (learning rate modulation, thresholds, etc) and some performance statistics use simple cascaded time-integrals of spike-driven Ca at the Neuron level, in the  $CaSpk$  variables. The initial Ca level from spiking is just multiplied by a gain factor:

- $SpikeG$  (8 or 12) = gain multiplier on spike for computing  $CaSpk$ : increasing this directly affects the magnitude of the trace values, learning rate in Target layers, and other factors that depend on  $CaSpk$  values: `RLRate`, `UpdtThr`. `Prjn.KinaseCa.SpikeG` provides an additional gain factor specific to the synapse-level trace factors, without affecting neuron-level  $CaSpk$  values. Larger networks require higher gain factors at the neuron level -- 12, vs 8 for smaller.

The cascaded integration of these variables is:

```
[] CaSpkM += (SpikeG * Spike - CaSpkM) / MTau  CaSpkP += (CaSpkM - CaSpkP) / PTau  CaSpkD
+= (CaSpkP - CaSpkD) / DTau
sendspike
```

### *SendSpike*

For each Neuron, if  $Spike \neq 0$ , then iterate over `SendPrjns` for that layer, and for each sending Synapse, `Prjn.GScale.Scale` (computed projection scaling, see `projection-scaling`Projection scaling) is multiplied by the synaptic weight  $Wt$ , and added into the `GBuf` buffer for each receiving neuron, at the ring index for `Com.Delay` (such that it will be added that many cycles later in `GFmSpikes`). The `PIBuf` for the inhibitory pool of each receiving neuron is also incremented.

This is expensive computationally because it requires traversing all of the synapses for each sending neuron, in a sparse manner due to the fact that few neurons are typically spiking at any given cycle.

```
sendsynca-recvsynca
```

### *SendSynCa, RecvSynCa*

If synapse-level calcium (Ca) is being used for the trace *Credit* assignment factor in learning, then two projection-level functions are called across all projections, which are optimized to first filter by any sending

neurons that have just spiked (`SendSynCa`) and then any receiving neurons that spiked (`RecvSynCa`) -- `Ca` only needs to be updated in these two cases. This major optimization is only possible when using the simplified purely spike-driven form of `Ca` as in the `CaSpk` vars above. Another optimization is to exclude any neurons for which `CaSpkP` and `CaSpkD` are below a low update threshold `UpdtThr = 0.01`.

After filtering, the basic cascaded integration shown above is performed on synapse-level variables where the immediate driving `Ca` value is the product of `CaSyn` on the `recv` and `send` neurons times a `SpikeG` gain factor:

- `CaM += (SpikeG * send.CaSyn * recv.CaSyn - CaM) / MTau`

The cascading is optimized to occur only at the time of spiking by looping over intervening time steps since last spike-driven update.

learning-dwt-wtfmdwt

*Learning: DWt, WtFmDWt*

dwt

*DWt*

After 200 cycles of neuron updating per above, the synaptic weight changes `DWt` are computed according to the trace-kinase learning algorithm:

- `DWt = Error * Credit`

The *Error* gradient component of this weight change was shown above, in terms of the receiving neuron's `CaP` (CaMKII LTP) and `CaD` (DAPK1 LTD) kinase activity:

- `Error = CaP - CaD`

The *Credit* assignment component is the *trace*, based on the longest time-scale cascaded synaptic `Ca` value, `CaD` as updated in the above functions:

- `Tr += (CaD - Tr) / Tau // Tau = 1 or 2+ trials`

Along with a `RLRate` factor that represents the derivative of the receiving activation, which is updated for each neuron at the end of the *plus* phase prior to doing `DWt`:

- `RLRate = CaSpkD * (Max - CaSpkD) * (ABS(CaSpkP - CaSpkD) / MAX(CaSpkP - CaSpkD))`

– `Max` = maximum `CaSpkD` value across the layer

Thus, the complete learning function is:

- `DWt = (recv.CaP - recv.CaD) * Tr * recv.RLRate`

The soft weight bounding is applied at the time of computing the DWt, as a function of the Linear weight value LWt (explained below in WtFmDWt) as follows:

```
[] if DWt > 0 { DWt *= (1 - LWt) } else { DWt *= LWt }
```

There are also alternative learning functions supported, that use neuron-level Ca instead of synaptic Ca (selected by the `NeuronCa` option) for the trace factor, and are thus significantly faster, but do not work as well in general, especially in large networks on challenging tasks.

```
wtfmdwt
```

### *WtFmDWt*

Synaptic weights Wt are typically updated after every weight change, but multiple DWt changes can be added up in a mini-batch (often when doing data-parallel learning across multiple processors). With the SWt and contrast enhancement required to compensate for the soft weight bounding (which was also a long-time part of the Leabra algorithm), there are *three* different weight values at each synapse:

- Wt = effective synaptic weight value, determining how much conductance one spike drives on the receiving neuron, representing the actual number of effective AMPA receptors in the synapse.  $Wt = SWt * WtSig(LWt)$ , where WtSig produces values between 0-2 based on LWt, centered on 1.
- LWt = rapidly learning, linear weight value -- learns according to the lrate specified in the connection spec. Biologically, this represents the internal biochemical processes that drive the trafficking of AMPA receptors in the synaptic density. Initially all LWt are .5, which gives 1 from WtSig function.
- SWt = slowly adapting structural weight value, which acts as a multiplicative scaling factor on synaptic efficacy: biologically represents the physical size and efficacy of the dendritic spine. SWt values adapt in an outer loop along with synaptic scaling, with constraints to prevent runaway positive feedback loops and maintain variance and further capacity to learn. Initial variance is all in SWt, with LWt set to .5, and scaling absorbs some of LWt into SWt.

First, the LWt is updated from the DWt:

- $LWt += DWt$

Then Wt is updated therefrom:

- $Wt = SWt * Sigmoid(LWt)$
- $Sigmoid(W) = 1 / (1 + ((1-W)/W)^6)$

Thus, the SWt provides a multiplicative constraint on the weights, and the LWt drives a more extreme, contrast-enhanced value on the weights, which counteracts the compression created by the soft weight bounding.

```
slowadapt-updates-target-activity-rescaling-sw
```

### *SlowAdapt Updates: Target Activity Rescaling, SWt*

Every `SlowInterval` (100) Trials, the `SlowAdapt` methods are called on all Layers and then Projections, which perform the following. These are essential constraints on learning that break the positive feedback loops while preserving effective error-driven learning.

target-vs-average-activity

### Target vs. Average Activity

First, when the network is initialized, a `TrgAvg` value is assigned to each neuron by uniformly sampling within a range of target values (0.5 - 2.0) and permuting the values among the set of neurons. This target is then updated as a function of the receiving unit error-gradient, subject to a zero-sum constraint across the relevant Pool of neurons:

- `DTrgAvg += ErrLRate * (CaSpkP - CaSpkD) // ErrLRate = .02`
- `TrgAvg += DTrgAvg - AVG(DTrgAvg) // zero-sum`

After every Trial, the neuron's actual average activation `ActAvg` is updated in a running-average manner:

- `ActAvg += (ActM - ActAvg) / LongAvgTau // LongAvgTau = 20 trials`

Then, in `SlowAdapt`, the `ActAvg` values are normalized into a proportion relative to the averages within the Pool a neuron belongs to, and the difference between this and the `TrgAvg` recorded:

- `AvgPct = ActAvg / SUM(ActAvg)`
- `AvgDif = AvgPct - TrgAvg`

This `AvgDif` value then drives synaptic rescaling per below.

swt-update

### SWt Update

The `SWt` is updated from `DSWt` which is accumulated from all the ensuing `DWt` values, with soft bounding applied and zero-sum:

```
[] if DSWt >= 0 { DSWt *= (SWt.Limit.Max - SWt) } else { DSWt *= (SWt - SWt.Limit.Min) } SWt
+= SWt.Adapt.LRate * (DSWt - AVG(DSWt) // AVG over Recv synapses per Prjn LWt = SigInverse(Wt /
SWt) // inverse of sigmoid
```

The learning rate here is typically slow, on the order 0.001 or even lower in large networks.

The updating of `LWt` in this way preserves the current `Wt` value despite changes in the `SWt` multiplier -- in effect the `LWt` absorbs the changes in `SWt` to preserve the current `Wt` value.

synaptic-rescaling

### Synaptic Rescaling

Finally, the `LWt` values are rescaled as a function of the `AvgDif` values reflecting the deviation in average activity relative to the target as computed above, using a soft-bounding update:

```
[] if AvgDif > 0 { LWt += SynScaleRate * (1 - LWt) * AvgDif * SWt } else { LWt += SynScaleRate *
LWt * AvgDif * SWt } Wt = SWt * Sigmoid(LWt)
```

This updates all the learned weights, and consequently the effective weights, moving in the direction to reduce the difference between the actual average activation and the target.

projection-scaling



## Projection scaling

The `Ge` and `Gi` synaptic conductances computed from a given projection from one layer to the next reflect the number of receptors currently open and capable of passing current, which is a function of the activity of the sending layer, and total number of synapses. We use a set of equations to automatically normalize (rescale) these factors across different projections, so that each projection has roughly an equal influence on the receiving neuron, by default.

The most important factor to be mindful of for this automatic rescaling process is the expected activity level in a given sending layer. This is set initially to `Layer.Inhib.ActAvg.Init`, and adapted from there by the various other parameters in that `Inhib.ActAvg` struct. It is a good idea in general to set that `Init` value to a reasonable estimate of the proportion of activity you expect in the layer, and in very small networks, it is typically much better to just set the `Fixed` flag and keep this `Init` value as such, as otherwise the automatically computed averages can fluctuate significantly and thus create corresponding changes in input scaling. The default `UseFirst` flag tries to avoid the dependence on the `Init` values but sometimes the first value may not be very representative, so it is better to set `Init` and turn off `UseFirst` for more reliable performance.

Furthermore, we add two tunable parameters that further scale the overall conductance received from a given projection (one in a *relative* way compared to other projections, and the other a simple *absolute* multiplicative scaling factor). These are some of the most important parameters to configure in the model -- in particular the strength of top-down "back" projections typically must be relatively weak compared to bottom-up forward projections (e.g., a relative scaling factor of 0.1 or 0.2 relative to the forward projections).

The scaling contributions of these two factors are:

- `GScale = WtScale.Abs * (WtScale.Rel / Sum(all WtScale.Rel))`

Thus, all the `Rel` factors contribute in proportion to their relative value compared to the sum of all such factors across all receiving projections into a layer, while `Abs` just multiplies directly.

In general, you want to adjust the `Rel` factors, to keep the total `Ge` and `Gi` levels relatively constant, while just shifting the relative contributions. In the relatively rare case where the overall `Ge` levels are too high or too low, you should adjust the `Abs` values to compensate.

Typically the `Ge` value should be between .5 and 1, to maintain a reasonably responsive neural response, and avoid numerical integration instabilities and saturation that can arise if the values get too high. You can record the `Layer.Pools[0].Inhib.Ge.Avg` and `.Max` values at the epoch level to see how these are looking -- this is especially important in large networks, and those with unusual, complex patterns of connectivity, where things might get out of whack.

automatic-rescaling

## Automatic Rescaling

Here are the relevant factors that are used to compute the automatic rescaling to take into account the expected activity level on the sending layer, and the number of connections in the projection. The actual code is in `axon/layer.go`: `GScaleFmAvgAct()` and `axon/act.go` `SLayActScale`

- `savg` = sending layer average activation
- `snu` = sending layer number of units
- `ncon` = number of connections
- `slayActN` = `int(Round(savg * snu))` -- must be at least 1

- `sc` = scaling factor, which is roughly  $1 / \text{expected number of active sending connections}$ .
- `if ncon == snu: -- full connectivity`
  - `sc = 1 / slayActN`
- `else: -- partial connectivity -- trickier`
  - `avgActN = int(Round(savg * ncon))` -- avg proportion of connections
  - `expActN = avgActN + 2` -- add an extra 2 variance around expected value
  - `maxActN = MIN(ncon, slayActN)` -- can't be more than number active
  - `expActN = MIN(expActN, maxActN)` -- constrain
  - `sc = 1 / expActN`

This `sc` factor multiplies the `GScale` factor as computed above.

important-stats

### Important Stats

The only way to manage the complexity of large spiking nets is to develop advanced statistics that reveal what is going on, especially when things go wrong. These include:

- Basic "activation health": proper function depends on neurons remaining in a sensitive range of excitatory and inhibitory inputs, so these are monitored. Each layer has `ActAvg` with `AvgMaxGeM` reporting average maximum minus-phase `Ge` values -- these are what is regulated relative to `Act.GTarg.GeMax`, but also must be examined early in training to ensure that initial excitation is not too weak. The layer `Inhib.ActAvg.Init` can be set to adjust -- and unlike in *Leabra*, there is a separate `Target` value that controls adaptation of layer-level inhibition.
- PCA of overall representational complexity. Even if individual neurons are not hogging the space, the overall complexity of representations can be reduced through a more distributed form of hogging. PCA provides a good measure of that.

appendix-specialized-bg-pfc-da-etc-algorithms

## Appendix: Specialized BG / PFC / DA / Etc Algorithms

There are various extensions to the algorithm that implement special neural mechanisms associated with the prefrontal cortex and basal ganglia `pbwmPBWM`, dopamine systems `pvlvPVLV`, the hippocampus `Hippocampus`, and predictive learning and temporal integration dynamics associated with the thalamocortical circuits `deepaxonDeepAxon`. All of these are (will be) implemented as additional modifications of the core, simple `axon` implementation, instead of having everything rolled into one giant hairball as in the original C++ implementation.

This repository contains specialized additions to the core algorithm described here:

- <https://github.com/emeryf/axon/blob/master/deepdeep> has the DeepAxon mechanisms for simulating the deep neocortical <-> thalamus pathways (wherein basic Axon represents purely superficial-layer processing)
- <https://github.com/emeryf/axon/blob/master/rlpbwm> has basic reinforcement learning models such as Rescorla-Wagner and TD (temporal differences).
- <https://github.com/emeryf/axon/blob/master/pbwm1pbwm> has the prefrontal-cortex basal ganglia working memory model (PBWM).
- <https://github.com/emeryf/axon/blob/master/hiphip> has the hippocampus specific learning mechanisms.

appendix-kinase-trace-learning-rule-derivation

## Appendix: Kinase-Trace Learning Rule Derivation

To begin, the original *GeneRec* references (O'Reilly, 1996) derivation of CHL (contrastive hebbian learning) from error backpropagation goes like this:

$$\frac{\partial E}{\partial w} = \frac{\partial E}{\partial y} \frac{\partial y}{\partial g} \frac{\partial g}{\partial w}$$

where  $E$  is overall error,  $w$  is the weight,  $y$  is recv unit activity,  $g$  is recv conductance (net input), and  $x$  is sending activity. For a simple neural network:

$$g = \sum x w$$

$$y = f(g)$$

This chain rule turns into:

$$dW = \frac{\partial E}{\partial w} = \left[ \left( \sum_i x_i^+ - \sum_i x_i^- \right) w \right] y' x = (g^+ - g^-) y' x$$

Thus, the *Error* factor is  $(g^+ - g^-)$  and  $y' x$  is the *Credit* factor. In words, the error signal is received by each unit in the form of their weighted net input from all other neurons -- the error is the temporal difference in this net input signal between the plus and minus phases. And the credit assignment factor is the sending unit activity  $x$  times the derivative of activation function.

The presence of this derivative is critical -- and has many tradeoffs embedded within it, as discussed later (e.g., the ReLU eliminates the derivative by using a mostly linear function, and thereby eliminates the *vanishing gradient* problem that otherwise occurs in sigmoidal activation functions).

The original GeneRec derivation of CHL mixes these factors by approximating the derivative of the activation function using the discrete difference in receiving activation state, such that:

$$(g^+ - g^-) y' \approx y^+ - y^-$$

In the GeneRec derivation, the approximate midpoint integration method, and symmetry preservation, cause the terms to get mixed together with the sending activations, producing the CHL algorithm.

To derive the new trace-enabling rule, we avoid this mixing, and explore learning using the more separable Error \* Credit form. In practice, the key issue is *on what variable is the temporal difference computed*: just using raw net input turns out to be too diffuse -- the units end up computing too similar of error gradients, and the credit assignment is not quite sufficient to separate them out.

In the Axon framework in particular, the weights are constrained to be positive, and especially at the start of learning, the net input terms are all fairly close in values across units. The lateral inhibition provides the critical differentiation so that only a subset of neurons are active, and thus having some contribution of the actual receiving activity is critical for a learning rule that ends up having different neurons specializing on

different aspects of the problem. The relative lack of this kind of differential receiver-based credit assignment in backprop nets is a critical difference from the CHL learning rule -- in the GeneRec derivation, it arises from making the learning rule symmetric, so that the credit assignment factor includes both sides of the synapse.

In short, backprop is at one end of a continuum where the only credit assignment factor is presynaptic activity, and existing weights provide a "filter" through which the Error term is processed. At the other end is the symmetric CHL equation where  $\text{pre} * \text{post} (xy)$  is the credit assignment factor in effect, and the "trace" equation is somewhere in between.

appendix-neural-data-for-parameters

### Appendix: Neural data for parameters

See <https://github.com/emeryf/axon/tree/master/chanschans> for good sources for many constants and equations.

- Brunel00: [https://nest-simulator.readthedocs.io/en/nest-2.20.1/auto\\_examples/brunel\\_alpha\\_nest.html](https://nest-simulator.readthedocs.io/en/nest-2.20.1/auto_examples/brunel_alpha_nest.html)
- SahHestrinNicoll90 -- actual data for AMPA, NMDA
- XiangHuguenardPrince98a -- actual data for GABAa

Axonal conduction delays:

- [http://www.scholarpedia.org/article/Axonal\\_conduction\\_delay](http://www.scholarpedia.org/article/Axonal_conduction_delay)
  - thalamocortical is very fast: 1.2 ms
  - corticocortical axonal conduction delays in monkey average 2.3 ms (.5 to 8 range)
  - corpus callosum (long distance) average around 10 ms
  - Brunel00: 1.5 ms

AMPA rise, decay times:

- SHN90: rise times 1-3ms -- recorded at soma -- reflect conductance delays
- SHN90: decay times 4-8ms mostly
- Brunel00: 0.5 ms -- too short!

GABAa rise, decay times:

- XHP98a: 0.5ms rise, 6-7ms decay

appendix-dendritic-dynamics

## Appendix: Dendritic Dynamics

A series of models published around the year 2000 investigated the role of active dendritic channels on signal integration across different dendritic compartments (Migliore et al, 1999; Poirazi et al, 2003; Jarsky et al, 2005) -- see references Spruston (2008), Poirazi & Papoutsi (2020) for reviews. A common conclusion was that the A-type K channel can be inactivated as a result of elevated  $V_m$  in the dendrite, driving a nonlinear gating-like interaction between dendritic inputs: when enough input comes in (e.g., from 2 different projections), then the rest of the inputs are all integrated more-or-less linearly, but below this critical threshold, inputs are much more damped by the active A-type K channels. There are also other complications associated with VGCC L-type and T-type voltage-gated Ca channels which can drive Ca spikes, to amplify regular AMPA conductances, relative weakness and attenuation of active HH Na spiking channels in dendrites, and issues of where inhibition comes in, etc. See following figure from Spruston (2008) for a summary of some key traces:

Here are some specific considerations and changes to capture some of these dynamics:

- The `Kdr` delayed rectifier channel, part of the classical HH model, resets the membrane potential back to resting after a spike -- according to detailed traces from the <https://github.com/ccnlab/kinase/tree/main/sims/urakubo> Urakubo et al., 2008 model with this channel, and the above figures, this is not quite an instantaneous process, with a time constant somewhere between 1-2 msec. This is not important for overall spiking behavior, but it is important when  $V_m$  is used for more realistic Ca-based learning (as in the Urakubo model). This more realistic  $V_m$  reset behavior is captured in Axon via the `RTau` time constant, which decays  $V_m$  back to  $V_{mR}$  within the `Tr` refractory period of 3 msec, which fits well with the Urakubo traces for isolated spikes.
- The `Dend` params specify a `GbarExp` parameter that applies a fraction of the `Exp` slope to  $V_{mDend}$ , and a `GbarR` param that injects a proportional amount of leak current during the spike reset (`Tr`) window to bring the  $V_m$  back down a bit, reflecting the weaker amount of `Kdr` out in the dendrites. This produces traces that resemble the above figure, as shown in the following run of the `examples/neuron` model, comparing  $V_{mDend}$  with  $V_m$ . Preliminary indications suggest this has a significant benefit on model performance overall (on `ra25` and `fsa` so far), presumably by engaging NMDA and GABAB channels better.
- As for the broader question of more coincidence-driven dynamics in the dendrites, or an AND-like mutual interdependence among inputs to different branches, driven by A-type K channels, it is likely that in the awake behaving context (*in activo*) as compared to the slices where these original studies were done, there is always a reasonable background level of synaptic input such that these channels are largely inactivated anyway. This corresponds to the important differences between upstate / downstate that also largely disappear in awake behaving vs. anesthetized or slice preps. Nevertheless, it is worth continuing to investigate this issue and explore the potential implications of these mechanisms in actual running models. TODO: create `atype` channels in `glong` (rename to something else, maybe just `chans` for channels)

TODO: GaoGrahamZhouEtAl20




Figure 1: The XCAL dWt function, showing direction and magnitude of synaptic weight changes (dWt) as a function of the short-term average activity of the sending neuron ( $x$ ) times the receiving neuron ( $y$ ). This quantity is a simple mathematical approximation to the level of postsynaptic  $Ca^{++}$ , reflecting the dependence of the NMDA channel on both sending and receiving neural activity. This function was extracted directly from the detailed biophysical Urakubo et al., (2008) model, by fitting a piecewise linear function to the synaptic weight change behavior that emerges from it as a function of a wide range of sending and receiving spiking patterns.

## Stabilizing Effects of NMDA and GABA-B

## Basic Supervised Category Learning in the MNIST Digit Recognition Task

## Large-scale 100-way Classification Learning of Realistic 3D Images

## Predictive Learning

## General Discussion

\*

## References

- Bi, G.-q., & Poo, M.-m. (1998, December). Synaptic modifications in cultured hippocampal neurons: Dependence on spike timing, synaptic strength, and postsynaptic cell type. *The Journal of Neuroscience*,

- 18(24), 10464–10472. Retrieved 2016-01-27, from <http://www.jneurosci.org/content/18/24/10464>
- Brette, R., & Gerstner, W. (2005, November). Adaptive exponential integrate-and-fire model as an effective description of neuronal activity. *Journal of Neurophysiology*, 94(5), 3637–3642. doi: 10.1152/jn.00686.2005
- Gerstner, W., & Naud, R. (2009, October). How Good Are Neuron Models? *Science*, 326(5951), 379–380. Retrieved from <http://www.sciencemag.org>
- Graupner, M., & Brunel, N. (2012, March). Calcium-based plasticity model explains sensitivity of synaptic changes to spike pattern, rate, and dendritic location. *Proceedings of the National Academy of Sciences*, 109(10), 3991–3996. doi: 10.1073/pnas.1109359109
- Kheradpisheh, S. R., Ganjtabesh, M., Thorpe, S. J., & Masquelier, T. (2018, March). STDP-based spiking deep convolutional neural networks for object recognition. *Neural Networks*, 99, 56–67. doi: 10.1016/j.neunet.2017.12.005
- LeCun, Y., Bengio, Y., & Hinton, G. (2015, May). Deep learning. *Nature*, 521(7553), 436–444. doi: 10.1038/nature14539
- Lillicrap, T. P., Santoro, A., Marris, L., Akerman, C. J., & Hinton, G. (2020, June). Backpropagation and the brain. *Nature Reviews Neuroscience*, 21(6), 335–346. doi: 10.1038/s41583-020-0277-3
- McKee, K. L., Crandell, I. C., Chaudhuri, R., & O'Reilly, R. C. (2021, November). Locally learned synaptic dropout for complete Bayesian inference. *arXiv:2111.09780 [q-bio, stat]*. Retrieved 2021-11-20, from <http://arxiv.org/abs/2111.09780>
- Miller, E. K., Erickson, C. A., & Desimone, R. (1996, August). Neural mechanisms of visual working memory in prefrontal cortex of the macaque. *Journal of Neuroscience*, 16(16), 5154–5167. Retrieved from <http://www.ncbi.nlm.nih.gov/pubmed/8756444>
- O'Reilly, R. C. (1996, January). Biologically plausible error-driven learning using local activation differences: The generalized recirculation algorithm. *Neural Computation*, 8(5), 895–938. doi: 10.1162/neco.1996.8.5.895
- O'Reilly, R. C., Russin, J. L., Zolfaghar, M., & Rohrlich, J. (2020, June). Deep Predictive Learning in Neocortex and Pulvinar. *arXiv:2006.14800 [q-bio]*. Retrieved 2020-06-30, from <http://arxiv.org/abs/2006.14800>
- Rumelhart, D. E., Hinton, G. E., & Williams, R. J. (1986, January). Learning representations by back-propagating errors. *Nature*, 323(9), 533–536.
- Sanders, H., Berends, M., Major, G., Goldman, M. S., & Lisman, J. E. (2013, January). NMDA and GABAB (KIR) Conductances: The “Perfect Couple” for Bistability. *Journal of Neuroscience*, 33(2), 424–429. doi: 10.1523/JNEUROSCI.1854-12.2013
- Scellier, B., & Bengio, Y. (2017, May). Equilibrium propagation: Bridging the gap between energy-based models and backpropagation. *Frontiers in Computational Neuroscience*, 11. doi: 10.3389/fncom.2017.00024
- Shouval, H. Z., Wang, S. S.-H., & Wittenberg, G. M. (2010). Spike timing dependent plasticity: A consequence of more fundamental learning rules. *Frontiers in Computational Neuroscience*, 4(19). Retrieved from <http://www.ncbi.nlm.nih.gov/pubmed/20725599>
- Thorpe, S., Delorme, A., & Van Rullen, R. (2001). Spike-based strategies for rapid processing. *Neural networks*, 14, 715–725. Retrieved from <http://www.ncbi.nlm.nih.gov/pubmed/11665765>
- Whittington, J. C. R., & Bogacz, R. (2019, March). Theories of error back-propagation in the brain. *Trends in Cognitive Sciences*, 23(3), 235–250. doi: 10.1016/j.tics.2018.12.005
- Xie, X., & Seung, H. S. (2003, February). Equivalence of backpropagation and Contrastive Hebbian Learning in a layered network. *Neural Computation*, 15(2), 441–454. Retrieved from <http://www.ncbi.nlm.nih.gov/pubmed/12590814>

